

# (BQ-BQ) Pipeline Project Documentation

## Table of Contents

1. Overview
  2. Architecture
  3. Tools Used
  4. Objective
  5. Implementations
  6. Highlights
  7. Future Enhancements
  8. Conclusion
- 

## 1. Overview

This project implements a **BigQuery-to-BigQuery (BQ-BQ) pipeline** designed to automate the process of transferring, transforming, and loading data from a **public BigQuery dataset** (bigquery-public-data.fhir\_synthea.claim) into a **target BigQuery dataset** (synthea\_ds.accident\_data) inside the project **ejazgcp**.

The pipeline is designed to:

- Copy data from the source dataset.
- Transform and load it into a partitioned and clustered table for efficient querying.
- Keep the target table up-to-date using a **MERGE operation**.

It uses **Apache Airflow** for orchestration and the **BigQueryInsertJobOperator** to execute queries.

---

## 2. Architecture

Architecture Flow:

1. **Source Dataset** – Public BigQuery dataset (bigquery-public-data.fhir\_synthea.claim).

2. **Target Dataset** – synthea\_ds in project ejazgcp.
3. **Target Table** – accident\_data partitioned by event\_timestamp and clustered by severity, accident\_type.
4. **Airflow DAG** – Executes MERGE logic for incremental updates.
5. **BigQuery** – Handles data transformations and load.

## Process Flow:

mathematica

CopyEdit

Public Dataset (claim table)



BigQuery Query Job (MERGE)



Target Dataset: synthea\_ds.accident\_data



Partitioned + Clustered Storage

---

## 3. Tools Used

- **Google BigQuery** – Data storage, querying, partitioning, clustering.
  - **Apache Airflow** – Workflow orchestration.
  - **BigQueryInsertJobOperator** – To submit SQL jobs to BigQuery from Airflow.
  - **SQL MERGE** – For incremental upserts.
  - **Python** – For defining Airflow DAG logic.
- 

## 4. Objective

- Create a **BQ-BQ data pipeline** to copy and transform claim data.
- Ensure the target table is **optimized for query performance** (partition + clustering).
- Maintain **up-to-date data** using incremental loading (MERGE).

- Schedule and manage the pipeline with Airflow.
- 

## 5. Implementations

### Step 1 – Project and Dataset Setup

- **Project Created:** ejazgcp
- **Target Dataset:** synthea\_ds
- **Target Table Schema:**

sql

CopyEdit

```
CREATE TABLE ejazgcp.synthea_ds.accident_data (  
    accident_id STRING,  
    patient_id STRING,  
    event_timestamp TIMESTAMP,  
    location_description STRING,  
    accident_type STRING,  
    involved_vehicles INT64,  
    injured_count INT64,  
    fatalities_count INT64,  
    hospital_admission BOOL,  
    reported_by STRING,  
    created_at TIMESTAMP,  
    updated_at TIMESTAMP  
)  
  
PARTITION BY DATE(event_timestamp)  
  
CLUSTER BY severity, accident_type;
```

---

### Step 2 – Initial Data Load

sql

CopyEdit

```
CREATE TABLE synthea_ds.accident_data AS  
SELECT * FROM `bigquery-public-data.fhir_synthea.claim`;
```

---

### Step 3 – Incremental Load Using MERGE

The MERGE operation:

- **Matches** existing rows on accident\_id.
- **Updates** if found.
- **Inserts** if not found.

sql

CopyEdit

```
MERGE `ejazgcp.synthea_ds.accident_data` T  
USING (  
    SELECT * FROM `bigquery-public-data.fhir_synthea.claim`  
) S  
ON T.accident_id = S.accident_id  
WHEN MATCHED THEN  
    UPDATE SET  
        patient_id = S.patient_id,  
        event_timestamp = S.event_timestamp,  
        location_description = S.location_description,  
        accident_type = S.accident_type,  
        involved_vehicles = S.involved_vehicles,  
        injured_count = S.injured_count,  
        fatalities_count = S.fatalities_count,  
        hospital_admission = S.hospital_admission,  
        reported_by = S.reported_by,  
        updated_at = CURRENT_TIMESTAMP
```

WHEN NOT MATCHED THEN

```
INSERT (  
    accident_id, patient_id, event_timestamp, location_description,  
    accident_type, involved_vehicles, injured_count, fatalities_count,  
    hospital_admission, reported_by, created_at, updated_at  
)  
VALUES (  
    S.accident_id, S.patient_id, S.event_timestamp, S.location_description,  
    S.accident_type, S.involved_vehicles, S.injured_count, S.fatalities_count,  
    S.hospital_admission, S.reported_by, CURRENT_TIMESTAMP, CURRENT_TIMESTAMP  
);
```

---

## Step 4 – Airflow DAG Implementation

python

CopyEdit

```
from airflow import DAG  
  
from airflow.providers.google.cloud.operators.bigquery import BigQueryInsertJobOperator  
  
from airflow.utils.dates import days_ago
```

```
default_args = {  
    'start_date': days_ago(1),  
    'retries': 1,  
}
```

```
with DAG(  
    dag_id='bq_to_bq_pipeline',  
    default_args=default_args,  
    schedule_interval=None,
```

```
catchup=False,  
tags=['bigquery', 'example'],  
) as dag:
```

```
bq_transfer = BigQueryInsertJobOperator(  
    task_id='transfer_claim_to_accident',  
    configuration={  
        "query": {  
            "query": """"<MERGE SQL ABOVE>""",  
            "useLegacySql": False  
        }  
    },  
    location='US'  
)
```

---

## 6. Highlights

- **Partitioned & Clustered Table** – Improves query speed and reduces cost.
- **Incremental Load** – MERGE ensures no duplicate rows and keeps data fresh.
- **Airflow Automation** – Easy scheduling and monitoring.
- **Scalability** – Handles large public dataset transfers efficiently.

---

## 7. Future Enhancements

- Add **Data Quality Checks** before loading.
- Integrate **Dataflow** for pre-BQ transformations.
- Add **Error Logging & Alerts** in Airflow.
- Implement **Automated Scheduling** for daily refresh.

---

## 8. Conclusion

This BQ-to-BQ pipeline ensures:

- **Automated, efficient, and scalable** data movement.
- **Optimized** table structure for analytics.
- **Maintainable** workflow using Airflow.